

Presented By Sophie A / 2026-01-21 / Everything Open 2026

Running Agentic AI Locally ➤➤➤

In this Workshop:

- What is **Open** vs **Closed** AI
- Installing the Tools
- **Reactive** vs **Agentic**
- Models for your Machine
- MCP Servers
- Advanced Tools



The screenshot shows a GitHub repository page for the project 'local-agentic-mcp-workshop'. The repository was created by 'starship-droid' and has a public status. It contains three files: 'README.md', 'Slides Running Agentic...', and 'meeting_minutes.md'. The 'README.md' file is the active tab, showing its content and a link to 'View Slides JanAI Quickstart'. Below the README, there's a section titled 'Getting Started' and 'Prerequisites for JanAI'. The repository has no releases or packages published.

starship-droid / local-agentic-mcp-workshop

Code Issues Pull requests Actions Projects Wiki Security

local-agentic-mcp-w... Public Pin Watch 0 Fork 0 Star 0

main Go to file + <> Code

starship-droid Slides added aa97c6b · yesterday

README.md Initial Commit yesterday

Slides Running Agentic... Slides added yesterday

meeting_minutes.md Initial Commit yesterday

README

Walkthrough: Running Agentic AI Locally on your Machine (GNU/Linux and Mac)

Github Repo View Slides JanAI Quickstart

Getting Started

Prerequisites for JanAI

Github Repo

<https://github.com/starship-droid/local-agentic-mcp-workshop/>

- Walkthrough README.md
- Slides PDF
- Example meeting_minutes.md

Download for your machine: jan.ai

The screenshot shows the homepage of the jan.ai website. At the top, there is a navigation bar with links to 'Docs', 'Changelog', 'Blog', and 'Handbook'. To the right of these links is a button labeled 'Download Jan' and social media icons for GitHub, X, LinkedIn, and a circular icon. Below the navigation bar, a banner displays a message: 'NEW ⭐ v0.7.5 is now live on GitHub. Check it out!'. The main title 'Ask Jan' is prominently displayed, featuring a yellow hand icon next to the word 'Ask'. Below the title, the tagline 'The best of open-source AI in an easy-to-use product.' is shown. Two calls to action are visible: a green button for 'Download for Mac' and a white button for 'Join community' with a user count of '15k+'. The background of the page features a cartoon illustration of a white, square-faced robot with a smiling mouth and pink cheeks, wearing a red cape, standing on a rocky ground with floating clouds. The overall theme is friendly and approachable.

Docs Changelog Blog Handbook

Download Jan

NEW ⭐ v0.7.5 is now live on GitHub. Check it out!

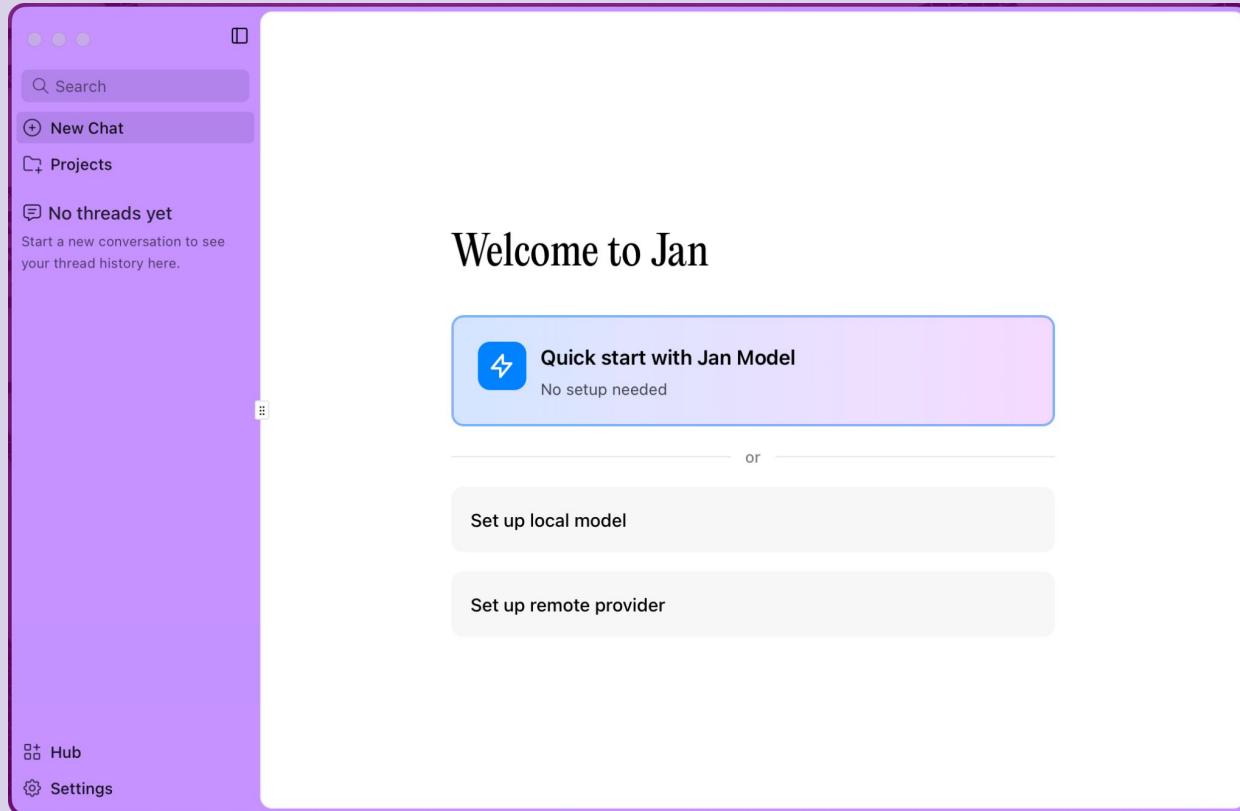
Ask Jan

The best of open-source AI in an easy-to-use product.

Download for Mac

Join community 15k+

+4.9M downloads



The image shows the home screen of the JanAI application. The interface has a light gray background with a purple header bar at the top. On the left side, there is a sidebar with a purple header containing three small circular icons, a search bar with a magnifying glass icon labeled "Search", a "New Chat" button with a plus sign and a speech bubble icon, and a "Projects" button with a folder icon. Below these are two sections: "No threads yet" with a speech bubble icon and "Start a new conversation to see your thread history here.", followed by a "Hub" button with a plus sign and a "Settings" button with a gear icon.

Welcome to Jan

Quick start with Jan Model
No setup needed

or

Set up local model

Set up remote provider

Home Screen of
JanAI

Hub

Search for models on Hugging Face...

Newest Downloaded

Jan-V1-4B-GGUF 2.3 GB [Download](#)

Jan-v1 is an advanced agentic language model built on the Qwen3-4B-thinking model, optimized for complex problem-solving and integration with...

By Janhq 86716 4 Show variants

Jan-V2-VL-High-Gguf 4.7 GB [Download](#)

Multimodal agent for stable long-horizon task execution in software environments.

By Janhq 54723 14 Show variants

Search

+ New Chat

Projects

Recents

Hello

...

...

Hub

Settings

☰ Download Qwen



5 - 15 min

The screenshot shows a user interface for downloading AI models. On the left, there's a sidebar with icons for New Chat, Projects, and No threads yet. The main area has a search bar with "Qwen" typed in, also highlighted by a red box. Below the search bar, three model cards are listed:

- Qwen3-VL-8B-Thinking-GGUF** (4.3 GB) - This card is partially visible at the top.
- Qwen3-VL-2B-Thinking-GGUF** (963.6 MB) - This card is highlighted with a red box. It describes the model as a vision-language model for image-text-to-text tasks with Unslloth optimizations. It was created by Unslloth, has 20309 downloads, 26 variants, and a "Show variants" toggle.
- Qwen3-4B-2507-Abliterated-GGUF** (2.3 GB) - This card is partially visible at the bottom.

Each model card includes a purple "Download" button and an information icon (i). A "Downloaded" toggle switch is located in the top right corner of the main area.

What is Local AI

- o
- o
- o
- o
- o
-

Vendor AI

Closed Weight

Used only online and through vendor tool and/or API

Unknown exactly how your data is used

Examples:

ChatGPT

Grok

Claude

DeepSeek (API / vendor-hosted)

Local AI

Open Weight / Open Source

Many options of how to run, eg locally through tool, through API,

Examples:

Llama

Mistral

Deepseek (Local / self-hosted)

How to Choose Models

Parameters		Reasoning	MCP	Vision
RAM	Model			
8GB	3B	Shows “thinking” on the screen Reasoning models can plan, loop, and achieve goals	Tools for LLM Examples: Searching on Web Filesystem Chrome Developer Tools	Model can analyse images Different models are better/worse at interpreting images
16GB	7B			
32GB	13B			

Compatibility

- Recommended for your device

Reactive vs Agentic AI

- o
- o
- o
- o
-

Reactive

Responds only to direct prompts

Single-step execution

No goal-setting or planning

Limited or no memory

Human controls every action

Agentic

Goal-driven and autonomous

Plans and executes multiple steps

Can decide what to do next

Uses memory and reasoning

Human supervises, not micromanages

Start Chat

The screenshot shows the Agentic AI application interface. On the left, there's a sidebar with icons for Search, New Chat (highlighted with a red box), Projects, No threads yet (with a note to start a new conversation), Hub, and Settings. The main area has a search bar at the top. A model card for "Qwen3-VL-8B-Thinking-1M-GGUF" is displayed, showing it's 4.3 GB in size and has been downloaded (indicated by a blue toggle switch with a red box around it). There's also a "Use" button (also with a red box) and some stats below: By Unslloth, 59334 views, 1 reply, and a magnifying glass icon. To the right, a chat window for "Jan" is open, showing a message "Hi, how are you?" and a response "How can I help you today?". Below the main window, two callout boxes point to the "Use" button and the message input field respectively. The callout for the "Use" button says "Shows model downloaded and usable". The callout for the message input field says "Tools available: image, attachment, web search".

Shows model downloaded and usable

Tools available: image, attachment, web search



MCP Servers





Settings

Settings

General

Attachments

Interface

Privacy

Model Providers >

Assistants

Shortcuts

Hardware

MCP Servers

Local API Server

HTTPS Proxy

Extensions

MCP Servers Experimental

Find more MCP servers at mcp.so

Allow All MCP Tool Permissions

When enabled, all MCP tool calls will be automatically approved without showing permission dialogs. This setting applies globally to all conversations, including new chats.

Tool call timeout (seconds)

Maximum time to wait for an MCP tool response before timing out.

30

Browser Use Proactive Mode Experimental

When enabled, automatically captures browser screenshots to provide additional context for browser automation tasks. Requires Jan Browser MCP with vision and tool use model.

● Jan Browser MCP Official

Transport: STDIO

Command: npx

Args: -y, search-mcp-server@latest

Env: BRIDGE_HOST=127.0.1, BRIDGE_PORT=*****

Requires Jan Browser Extension to be installed in your Chrome-based browser.

[Install Extension →](#)

● Browserman

Edit Filesystem MCP

The screenshot shows the Agentic AI application's settings interface. The left sidebar includes a search bar, 'New Chat' and 'Projects' buttons, a message 'No threads yet', and a note to start a new conversation. The 'MCP Servers' option is selected in the sidebar menu. The main content area displays four configuration sections: 'Fetch', 'Filesystem' (which has a red box around its edit icon), 'Sequential-Thinking', and 'Serper'. Each section shows transport (STDIO), command (uvx, npx, or npx), and arguments.

Settings

General
Attachments
Interface
Privacy
Model Providers >
Assistants
Shortcuts
Hardware
MCP Servers
Local API Server
HTTPS Proxy
Extensions

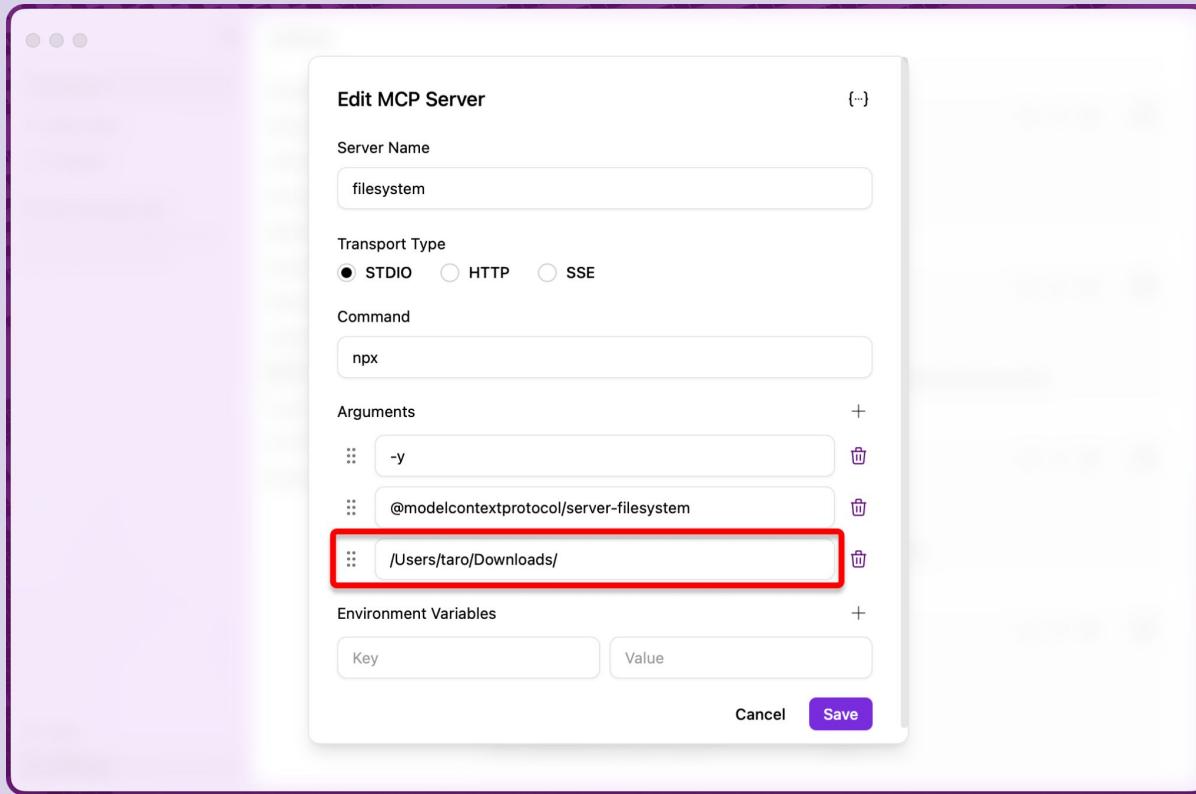
Fetch
Transport: STDIO
Command: uvx
Args: mcp-server-fetch

Filesystem (Edit icon highlighted)
Transport: STDIO
Command: npx
Args: -y, @modelcontextprotocol/server-filesystem, /path/to/other/allowed/dir

Sequential-Thinking
Transport: STDIO
Command: npx
Args: -y, @modelcontextprotocol/server-sequential-thinking

Serper
Transport: STDIO
Command: npx
Args: -y, serper-search-scrape-mcp-server
Env: SERPER_API_KEY=YOUR*****HERE

Enter Allowed Filepath



Enable Filesystem MCP

The screenshot shows the Agentic AI application's settings page. On the left, there's a sidebar with a search bar, new chat, projects, and a message about no threads yet. The main area has a title "Settings" and a sidebar with options: General, Attachments, Interface, Privacy, Model Providers (with a dropdown arrow), Assistants, Shortcuts, Hardware, MCP Servers (which is selected and highlighted in grey), Local API Server, HTTPS Proxy, and Extensions. At the bottom are "Hub" and "Settings" buttons.

MCP Servers

- Fetch**
Transport: STDIO
Command: uvx
Args: mcp-server-fetch
- Filesystem** (This item has a red box around its toggle switch)
Transport: STDIO
Command: npx
Args: -y, @modelcontextprotocol/server-filesystem, /path/to/other/allowed/dir
- Sequential-Thinking**
Transport: STDIO
Command: npx
Args: -y, @modelcontextprotocol/server-sequential-thinking
- Serpser**
Transport: STDIO
Command: npx
Args: -y, serper-search-scrape-mcp-server
Env: SERPER_API_KEY=YOUR*****HERE

Invalid path
causes error

An error dialog box is displayed. It has a title "Error" and a message "Something went wrong". Below it is a "Details" section with a collapsed arrowhead. The expanded details text reads: "Failed to start MCP server filesystem: Resolving dependencies Resolved, downloaded and extracted [2] Saved lockfile Error accessing directory /path/to/other/allowed/dir: ENOENT: no such file or directory, stat '/path/to/other/allowed/dir' path: '/path/to/other/allowed/dir', syscall: 'stat', errno: -2, code: 'ENOENT'".

Please check the parameters according to the tutorial.

Cancel Copy



Edit MCP

The screenshot shows the Agentic AI MCP interface. At the top, there's a header bar with three dots on the left, the title "Edit MCP" in large letters, and a search bar with a magnifying glass icon. Below the header is a sidebar with a purple background containing icons for "Search", "New Chat" (highlighted in pink), "Projects", and "No threads yet". A message "Server filesystem activated successfully" is displayed in a black bar at the top right. The main area features a white card with a purple border. Inside the card, the text "Hi, how are you?" is displayed in large font, followed by "How can I help you today?". Below this is a text input field with placeholder text "Ask me anything...". To the right of the input field is a toolbar with icons for a file, settings, a green dot, a photo, a link, and a globe. The icon for the globe is highlighted with a red box. At the bottom of the card, there are two buttons: "Hub" and "Settings".

○
○
○
○
○

```
# Sprint 24 Planning & Cross-Team Sync – Meeting Minutes

**Date:** January 12, 2026
**Time:** 10:00-11:00 AM
**Facilitator:** Emma Rodriguez (Engineering Manager)

## Attendees
- Emma Rodriguez – Engineering Manager
- Alex Chen – Frontend Lead
- Priya Patel – Backend Lead
- Jordan Lee – Data Engineering
- Sofia Martinez – Data Science
- Michael Brown – UI/UX Designer
- Rachel Kim – QA Lead
- Tom Wilson – DevOps

## Agenda
1. Sprint 24 goals
2. Feature progress updates
3. Data pipeline and analytics updates
4. UI/UX review feedback
5. Risks and blockers

## Discussion Notes

### 1. Sprint 24 Goals (All Teams)
The main goal for Sprint 24 is to deliver the first beta of the user activity dashboard logic. Teams are expected to coordinate closely due to shared dependencies.

### 2. Frontend & UI/UX Updates
The frontend team has completed the initial React components for the dashboard layout. The UI/UX design currently assumes both daily and weekly aggregation views, and c

### 3. Backend & Data Pipeline Updates
Backend APIs for user activity metrics are mostly implemented but depend on changes. The data engineering team is refactoring ETL jobs to support incremental updates. Some metric definitions, such as "active user," need alignment between data science and product teams.

### 4. Analytics, Experiments, and Logging
Plans were discussed to add additional event logging to support future A/B testing. The data warehouse schema will be updated to store raw events alongside aggregated data for faster analysis.
```

Summarise Meeting

1. Download from Github to your MCP allowed directory:
meeting_mintues.md
2. Sample prompt to create new file with summary

Customise JanAI

Running Agentic AI Locally on your Machine - Sophie A

The screenshot shows the JanAI application interface. On the left, there's a sidebar with a search bar, 'New Chat' and 'Projects' buttons, and a message 'No threads yet'. Below these are buttons for 'Start a new conversation' and 'your thread history here.'. At the bottom of the sidebar are 'Hub' and 'Settings' buttons, with 'Settings' highlighted by a red box. The main area is titled 'Assistants' and contains a list of configuration options: General, Attachments, Interface, Privacy, Model Providers, Assistants (which is selected and highlighted by a red box), Shortcuts, Hardware, MCP Servers, Local API Server, HTTPS Proxy, and Extensions. To the right, a card for 'Jan' is displayed, showing its description: 'Jan is a helpful desktop assistant that can reason through complex tasks and use tools to complete them on the user's behalf.' It includes an edit button (highlighted by a red box) and a delete button.

Edit Assistant

Emoji Name: Jan

Description (optional): Jan is a helpful desktop assistant that can reason through complex tasks and use tools to complete them on the user's behalf.

Instructions

You are a helpful AI assistant. Your primary goal is to assist users with their questions and tasks to the best of your abilities.

When responding:

- Answer directly from your knowledge when you can
- Be concise, clear, and helpful
- Admit when you're unsure rather than making things up

If tools are available to you:

- Only use tools when they add real value to your response
- Use tools when the user explicitly asks (e.g., "search for...", "calculate...", "run this code")
- Use tools for information you don't know or that needs verification
- Never use tools just because they're available

When using tools:

- Use one tool at a time and wait for results

MCP Servers Everywhere



Make
your own!

Chrome Dev Tools	GitHub	AWS	Reference
Inspect DOM, CSS, and console logs Access network requests and performance traces Debug web pages programmatically through an agent	Browse repos, issues, PRs, discussions Automate repository tasks and summaries Extract code context and change history for reasoning	Query cloud resources (S3, EC2, etc.) Manage infrastructure and retrieve logs Enable cloud-aware agent workflows	Official MCP reference implementations for tools like Git, Memory, Fetch, and Time <u>https://github.com/modelcontext/protocol/servers</u>



Additional Tools



Local API

An OpenAI compatible API running locally.
Existing tools and code can use Jan instead of the ChatGPT API.

HTTPS Proxy

A local HTTPS proxy that intercepts model requests when the API endpoint can't be changed.

The screenshot shows the Agentic AI application interface with two floating windows displaying configuration options.

Top Window (Local API Server):

- Left sidebar: New Chat, Projects, No threads yet, Model Providers (selected), Assistants, Shortcuts, Hardware, MCP Servers, Local API Server (highlighted with a red box), HTTPS Proxy, Extensions.
- Right pane: Local API Server (Run an OpenAI-compatible server locally), Start Server button, Auto start (Automatically start the Local API Server when the application launches. Uses last used model, or picks the first available model if unavailable), Server Logs (Open Logs), API Documentation (Open Docs).

Bottom Window (HTTPS Proxy):

- Left sidebar: New Chat, Projects, No threads yet, Model Providers (selected), Assistants, Shortcuts, Hardware, MCP Servers, Local API Server, HTTPS Proxy (highlighted with a red box), Extensions.
- Right pane: Proxy (Proxy URL: http://proxy.example.com:8080), Authentication (Username, Password), No Proxy (localhost, 127.0.0.1, local), Ignore SSL Certificates (Allow self-signed or unverified certificates. This may be required for some proxies but reduces security. Only enable if you trust your proxy).

Continue the discussion



Linux & AI



Luma Linux Victoria

Thank you!

